Feather PDF

A fast PDF reader that stays light on memory.

This sample document is part of the Feather PDF repository. Open it to try the reader: the Contents sidebar on the left lists the sections below, the toolbar has page navigation, zoom and search, and you can select and copy any of this text.

Feather PDF is written in Rust on top of PDFium, the PDF engine used by Google Chrome. The interface is drawn with plain Win32 and GDI, so there is no browser engine or UI framework running in the background.

1. Reading

Scroll with the mouse wheel, the arrow keys, Page Up and Page Down, or the thin scrollbar on the right. Type a page number into the page box in the toolbar to jump straight to it.

Zoom with Ctrl and the mouse wheel; the point under the pointer stays where it is. Ctrl+0 fits the page width, Ctrl+2 fits the whole page, and Ctrl+1 shows the page at its real size. In fit-page mode, Page Down moves exactly one page.

Press F11 for full screen. The toolbar and sidebar step aside, and a small label shows the page number as you read. Esc brings everything back.

When you close a document, Feather PDF remembers the page, zoom and fit mode, and opens it there next time.

2. Finding and copying

Press Ctrl+F to open the search field in the toolbar. Enter jumps to the next match and Shift+Enter to the previous one. Every match on screen is marked in yellow, and the current one in orange.

Long documents are searched in short slices, so the window keeps responding while the search runs through hundreds of pages.

Drag across text to select it, double-click to select a word, or press Ctrl+A to select the whole page. Ctrl+C copies the selection. Words that were hyphenated across a line break are joined again when copied.

3. How it stays light

Only the pages on screen are rendered, at screen resolution. Rendered pages are kept in a cache with a fixed budget of 32 MB; the least recently used ones are dropped first, and the cache is emptied when the window is minimized.

At high zoom, only the visible part of a page is rendered, so a page shown at 800 percent does not need a bitmap hundreds of megabytes large.

The file itself is read on demand instead of being loaded into memory, and when the PDF engine builds up too much internal data from pages it has visited, the document is quietly reopened to release it.